



MongoDB

Open-source, High-performance,
Schema-free, Document-Oriented Database

Who Use MongoDB

Justin.tv



Boxed Ice

sourceforge



github
SOCIAL CODING

BUSINESS
INSIDER



MongoDB & CAP Principle

Visual Guide to NoSQL Systems



MongoDB features

- Collection storage;
- Dynamic query;
- Complete index of support;
- Query monitoring, Query optimization;
- Replication automatic failover;
- Support binary data and large objects;
- Auto-sharding Support cloud level of flexibility;


RDBMS

- Great for many applications
- Shortcomings
- Scalability
- Flexibility

Scalability & performance

**Front
User**

•Memcached



•Key/value
Stores



•MongoDB



•RDBMS

**Back
end**

JSON-style Documents

Example

```
{ title : "My First Post", author: "javablogger",  
  comments : [{ by: "Abe", text: "First" },  
               { by : "Ada", text : "Good post" } ]  
}
```

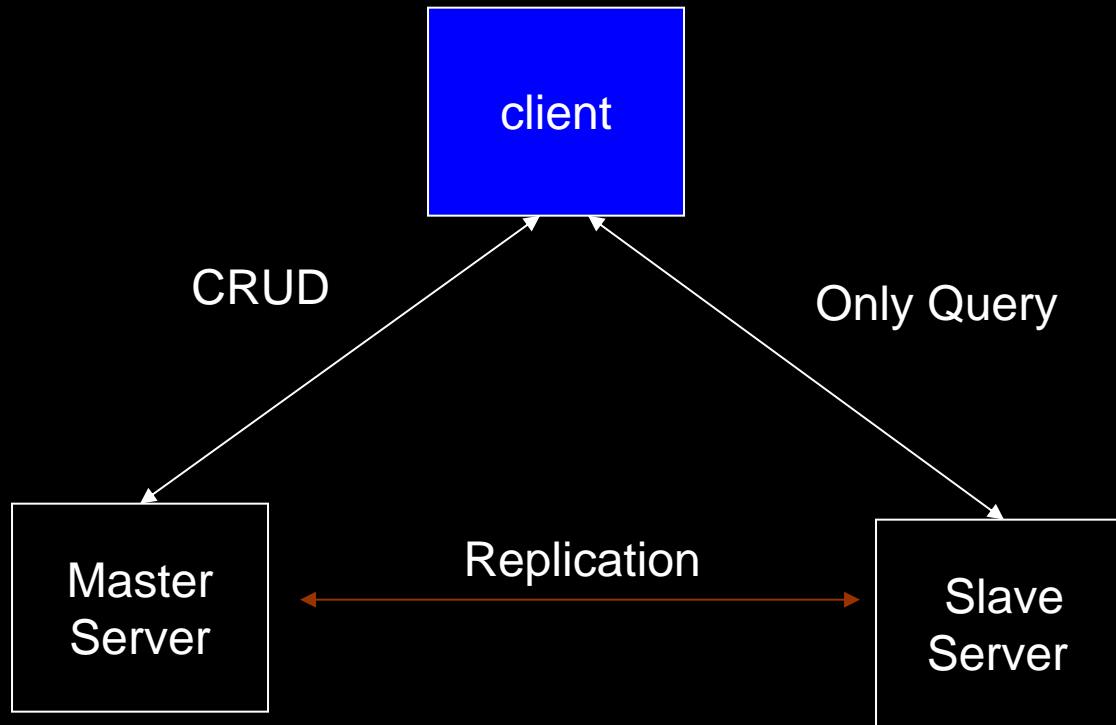
Schema-free

- Loosening constraints - added flexibility
- Dynamically typed languages (like Ruby!)
- Migrations

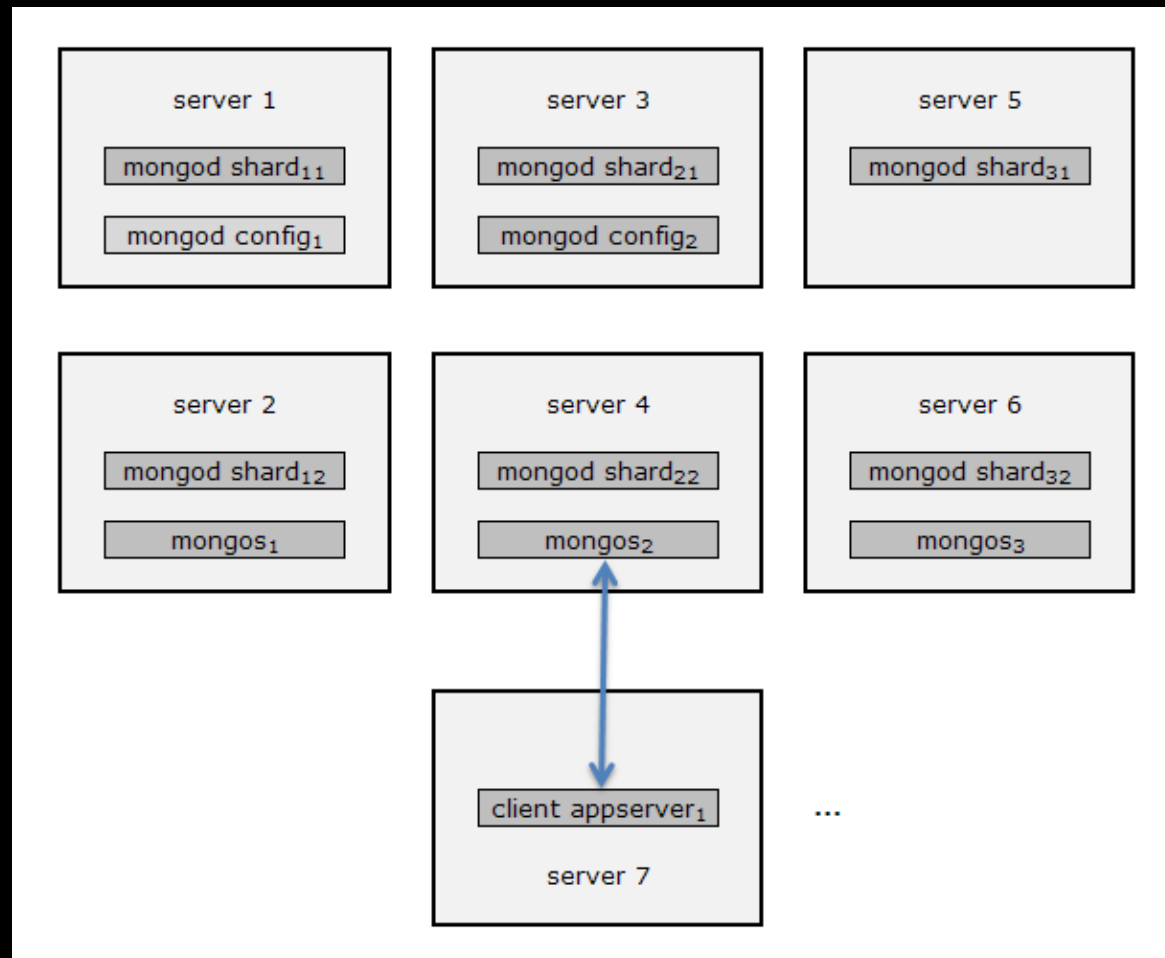
Dynamic queries

- Administration
- Ease of development
- Familiarity

Replication



Auto-sharding



MapReduce

Command syntax:

```
db.runCommand(  
  { mapreduce : <collection>,  
    map : <mapfunction>,  
    reduce : <reducefunction>  
    [, query : <query filter object>  
    [, sort : <sort the query. useful for optimization>  
    [, limit : <number of objects to return from collection>  
    [, out : <output-collection name>  
    [, keeptemp: <true|false>  
    [, finalize : <finalizefunction>  
    [, scope : <object where fields go into javascript global scope >  
    [, verbose : true] } );
```

Many Supported

- Platforms

Windows

Linux

Unix

bsd

- Languages

Java

C++

Ruby

PHP

- Docs

Administrative

Developer

Online API

Good at

- The web
- Caching
- High volume data
- Scalability

Less good at

- Highly transactional
- Ad-hoc business intelligence
- Problems that require SQL

MongoDB Basics

Document

- Unit of storage (think row)
- BSON (Binary JSON)
- Represented as a Hash

Collection

- Schema-free equivalent of a table
- Logical groups of documents
- Indexes are per-collection

`_id`

- Special key
- Present in all documents
- Unique across a Collection
- Any type you want

Blog back-end

Post

```
{:author => "mike",  
 :date => Time.new,  
 :text => "my blog post",  
 :tags => ["mongodb", "ruby"]}
```

Comment

```
{:author => "eliot",  
 :date => Time.new,  
 :text => "great post!"}
```

New post

```
post = {:author => "mike",  
        :date => Time.new,  
        :text => "my blog post",  
        :tags => ["mongodb", "ruby"]}
```

```
db["posts"].save(post)
```

Embedding a comment

```
c = {:author => "eliot",  
      :date => Time.new,  
      :text => "great post!"}
```

```
db["posts"].update({:_id => post[:_id]},  
                   {:$push => {:comments => c}})
```


Posts by author

```
db["posts"].find(:author => "mike")
```

Last 10 posts

```
db["posts"].find.sort([[:date, :desc]])  
.limit(10)
```

Posts in the last week

```
last_week = Time.utc(2009, 11, 12)  
db["posts"].find(:date => {:$gt => last_week})
```

Posts ending with
'Ruby'

```
db["posts"].find(:text => /Ruby$/)
```

Posts with a tag

```
db["posts"].find(:tags => "mongodb")
```

... and fast

```
db["posts"].create_index("tags")
```

Counting posts

```
db["posts"].count
```

```
db["posts"].find(:author => "mike").count
```

Basic paging

```
page = 2
```

```
page_size = 15
```

```
db["posts"].find.limit(page_size)  
                    .skip(page * page_size)
```

Migration: adding titles

- Easy - just start adding them:

```
post = {:author => "mike",  
       :date => Time.new,  
       :text => "another blog post",  
       :tags => ["RubyConf"],  
       :title => "Review from RubyConf"}
```

```
post_id = db["posts"].save(post)
```


Advanced queries

- \$gt, \$lt, \$gte, \$lte, \$ne, \$all, \$in, \$nin
- \$where

```
db["posts"].find :$where => "this.author == 'mike' ||  
                             this.title == 'hello'")
```

MongoMapper, Mongoid,
MongoRecord, etc.



MongoMapper

```
class User
  include MongoMapper::Document
  many :posts
end
```

```
class Post
  include MongoMapper::Document
  key :user_id, String
  key :title, String
end
```

```
user = User.create
user.posts.create(:title => 'Foo')
```

```
# would return post we just created
user.posts.find_by_title('Foo')
```

Other cool stuff

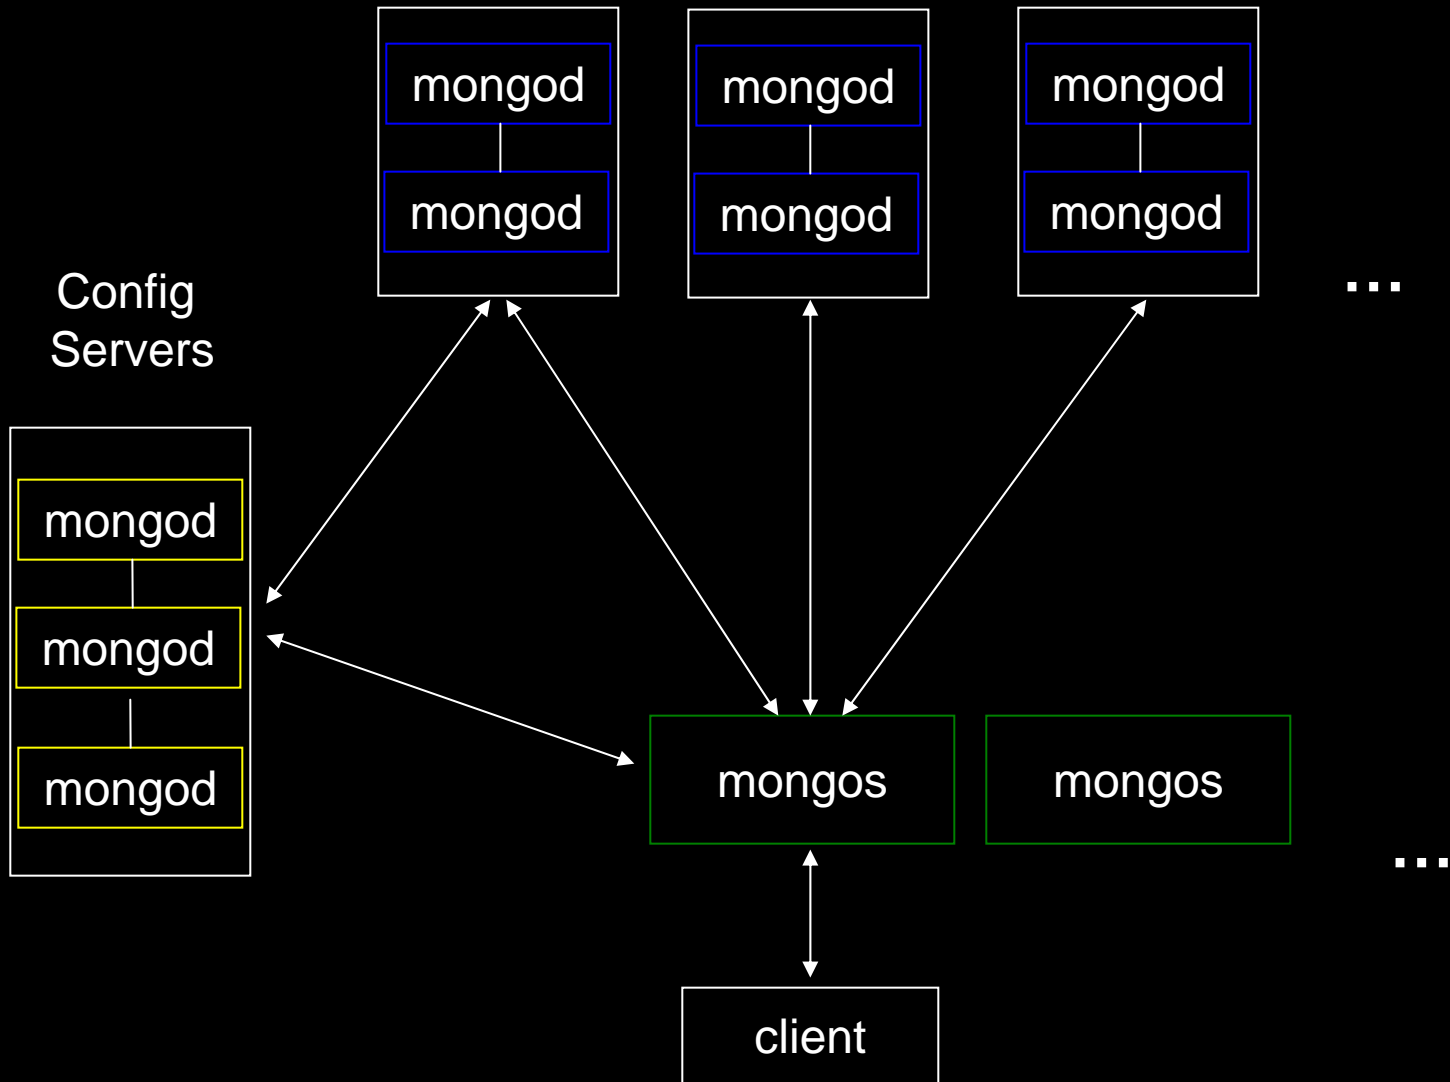
- Aggregation and map reduce
- Capped collections
- Unique indexes
- Mongo shell
- GridFS

Sharding

Terminology

- Shard key
- Chunk
 - Range of the value space
 - (collection, key, min_val, max_val)
- Shard
 - Single node (or replica pair)
 - Responsible for set of chunks

Shards



- **Download MongoDB**
<http://www.mongodb.org>
- **Try it out**
- **Let us know what you think!**

- <http://www.mongodb.org>
- <http://www.javabloger.com>
- njthnet@javabloger.com